Merging Data Sets

Data Wrangling in R

The `merge` function exists, don't use it

Joining in `dplyr`

- Merging/joining data sets together - usually on key variables, usually "id"
- `?join` - see different types of joining for `dplyr`
- Let's look at https://www.rstudio.com/wp-content/uploads/2015/02/data-wrangling-cheatsheet.pdf
- `inner_join(x, y)` - only rows that match for `x` and `y` are kept
- `full_join(x, y)` - all rows of `x` and `y` are kept
- `left_join(x, y)` - all rows of `x` are kept even if not merged with `y`
- `right_join(x, y)` - all rows of `y` are kept even if not merged with `x`
- `anti_join(x, y)` - all rows from `x` not in `y` keeping just columns from `x`.

Merging: Simple Data

`base` has baseline data for ids `1` to `10` and Age

```
base <- tibble(id = 1:10, Age = seq(55,60, length=10))
head(base, 2)
```

```
# A tibble: 2 x 2
     id    Age
  <int>  <dbl>
1     1     55
2     2   55.6
```

`visits` has ids 1 to 8, then 11 (new id), and 3 visits and outcome

```
visits <- tibble(id = c(rep(1:8, 3), 11), visit= c(rep(1:3, 8), 3),
                    Outcome = seq(10,50, length=25))
tail(visits, 2)
```

```
# A tibble: 2 x 3
     id visit Outcome
  <dbl> <dbl>   <dbl>
1     8     3    48.3
2    11     3      50
```

Inner Join

```
ij = inner_join(base, visits)
```

```
Joining, by = "id"
```

```
dim(ij)
```

```
[1] 24  4
```

```
tail(ij)
```

```
# A tibble: 6 x 4
     id   Age visit Outcome
  <dbl> <dbl> <dbl>   <dbl>
1     7  58.3     1      20
2     7  58.3     3    33.3
3     7  58.3     2    46.7
4     8  58.9     2    21.7
5     8  58.9     1      35
6     8  58.9     3    48.3
```

Left Join

```
lj = left_join(base, visits)
```

```
Joining, by = "id"
```

```
dim(lj)
```

```
[1] 26  4
```

```
tail(lj)
```

```
# A tibble: 6 x 4
     id   Age visit Outcome
  <dbl> <dbl> <dbl>   <dbl>
1     7  58.3     2    46.7
2     8  58.9     2    21.7
3     8  58.9     1    35
4     8  58.9     3    48.3
5     9  59.4    NA    NA
6    10  60      NA    NA
```

Right Join

```
rj = right_join(base, visits)

Joining, by = "id"

tail(rj, 3)
```

```
# A tibble: 3 x 4
     id    Age visit Outcome
  <dbl> <dbl> <dbl>   <dbl>
1     8  58.9     1      35
2     8  58.9     3    48.3
3    11  NA       3      50
```

Right Join: Switching arguments

```
rj2 = right_join(visits, base)

Joining, by = "id"

tail(rj2, 3)

# A tibble: 3 x 4
     id visit Outcome   Age
  <dbl> <dbl>   <dbl> <dbl>
1     8     3    48.3  58.9
2     9    NA      NA  59.4
3    10    NA      NA  60

identical(rj2, lj) ## after some rearranging

[1] TRUE
```

Full Join

```
fj = full_join(base, visits)

Joining, by = "id"

tail(fj, 4)

# A tibble: 4 x 4
     id   Age visit Outcome
  <dbl> <dbl> <dbl>   <dbl>
1     8  58.9     3    48.3
2     9  59.4    NA      NA
3    10  60      NA      NA
4    11  NA       3      50
```

Logging the joins

The `tidylog` package can show you log outputs from `dplyr` (newly added). You will need to install to use.

```
library(tidylog)
left_join(base, visits)
```

```
Joining, by = "id"

left_join: added 2 columns (visit, Outcome)

            > rows only in x     2

            > rows only in y  ( 1)

            > matched rows      24     (includes duplicates)

            >                       ====

            > rows total        26

# A tibble: 26 x 4
      id   Age visit Outcome
   <dbl> <dbl> <dbl>   <dbl>
 1     1    55     1      10
 2     1    55     3    23.3
 3     1    55     2    36.7
```

Using the `by` argument

By default - uses intersection of column names. If `by` specified, then uses that, but if other columns with same name, adds `suffix`.

```
base = base %>% mutate(x = 5)
visits = visits %>% mutate(x = 4)
head(full_join(base, visits))
```

```
Joining, by = c("id", "x")

# A tibble: 6 x 5
      id   Age     x visit Outcome
   <dbl> <dbl> <dbl> <dbl>   <dbl>
1      1  55       5    NA      NA
2      2  55.6     5    NA      NA
3      3  56.1     5    NA      NA
4      4  56.7     5    NA      NA
5      5  57.2     5    NA      NA
6      6  57.8     5    NA      NA
```

Using the **by** argument

```
head(full_join(base, visits, by = "id"))
```

```
# A tibble: 6 x 6
     id   Age   x.x visit Outcome   x.y
  <dbl> <dbl> <dbl> <dbl>   <dbl> <dbl>
1     1    55     5     1      10     4
2     1    55     5     3    23.3     4
3     1    55     5     2    36.7     4
4     2  55.6     5     2    11.7     4
5     2  55.6     5     1      25     4
6     2  55.6     5     3    38.3     4
```

```
head(full_join(base, visits, by = "id", suffix = c("_base", "_visit")))
```

```
# A tibble: 6 x 6
     id   Age x_base visit Outcome x_visit
  <dbl> <dbl>  <dbl> <dbl>   <dbl>   <dbl>
1     1    55      5     1      10       4
2     1    55      5     3    23.3       4
3     1    55      5     2    36.7       4
4     2  55.6      5     2    11.7       4
5     2  55.6      5     1      25       4
6     2  55.6      5     3    38.3       4
```

Using the `by` argument if column names different

```
base = base %>%
  select(-x) %>%
  mutate(myvar = 4)
visits = visits %>%
  select(-x) %>%
  mutate(MyVar = 4)
full_join(base, visits, by = c("id", "myvar" = "MyVar"))

# A tibble: 27 x 5
      id   Age myvar visit Outcome
   <dbl> <dbl> <dbl> <dbl>   <dbl>
 1     1  55       4     1    10
 2     1  55       4     3    23.3
 3     1  55       4     2    36.7
 4     2  55.6     4     2    11.7
 5     2  55.6     4     1    25
 6     2  55.6     4     3    38.3
 7     3  56.1     4     3    13.3
 8     3  56.1     4     2    26.7
 9     3  56.1     4     1    40
10     4  56.7     4     1    15
# … with 17 more rows
```